Title: Image Classification with Transfer Learning using Keras

Subtitle: Leveraging Pre-trained Models for Improved Accuracy

Swethadevi Ravipati (22082165)

Rani Panneru (22077318)

Muhammad Usman (22093856)

Hari Bahadur Gharti Magar (22075765)

Yogesh Pandit (22095146)

Geetha Babu (22073072)

Colab Link: https://colab.research.google.com/drive/1fC2G6J99dYpqkoXoHKfQISgwaAXxVkp4?usp=sharing

Problem Statement: The task is to develop an image classification system capable of accurately categorizing natural scenes from around the world. The dataset comprises images representing six categories: street, forest, sea, buildings, mountain, and glacier. The goal is to build a model that can effectively differentiate between these diverse environmental settings.

Objective:
Develop a deep learning-based image classification system using transfer learning techniques to accurately classify natural scene images into their respective categories. The model should achieve high accuracy while addressing the challenges posed by limited labeled data, diverse image characteristics, and imbalanced class distributions.

Evaluation Criteria:
The performance of the image classification model will be evaluated based on metrics such as accuracy, precision, recall, and F1-score. Additionally, the computational efficiency and scalability of the model will be considered to ensure its practical applicability in real-world scenarios.

Challenges:

- Limited Labeled Data: The dataset contains a limited number of labeled images for each category, making it challenging to train a robust classifier.
- Diverse Image Characteristics: Natural scenes exhibit a wide range of visual characteristics such as textures, colors, and structural elements, making it difficult for the model to generalize effectively.
- Imbalanced Classes: The distribution of images across different categories may be uneven, leading to imbalanced class distributions that could bias the model's learning.
- Computational Resources: Training deep neural networks requires significant computational resources, including GPU acceleration, which may not be readily available for all practitioners.
- Model Complexity: Selecting an appropriate model architecture and optimizing hyperparameters to balance model complexity and performance is crucial for achieving good results.

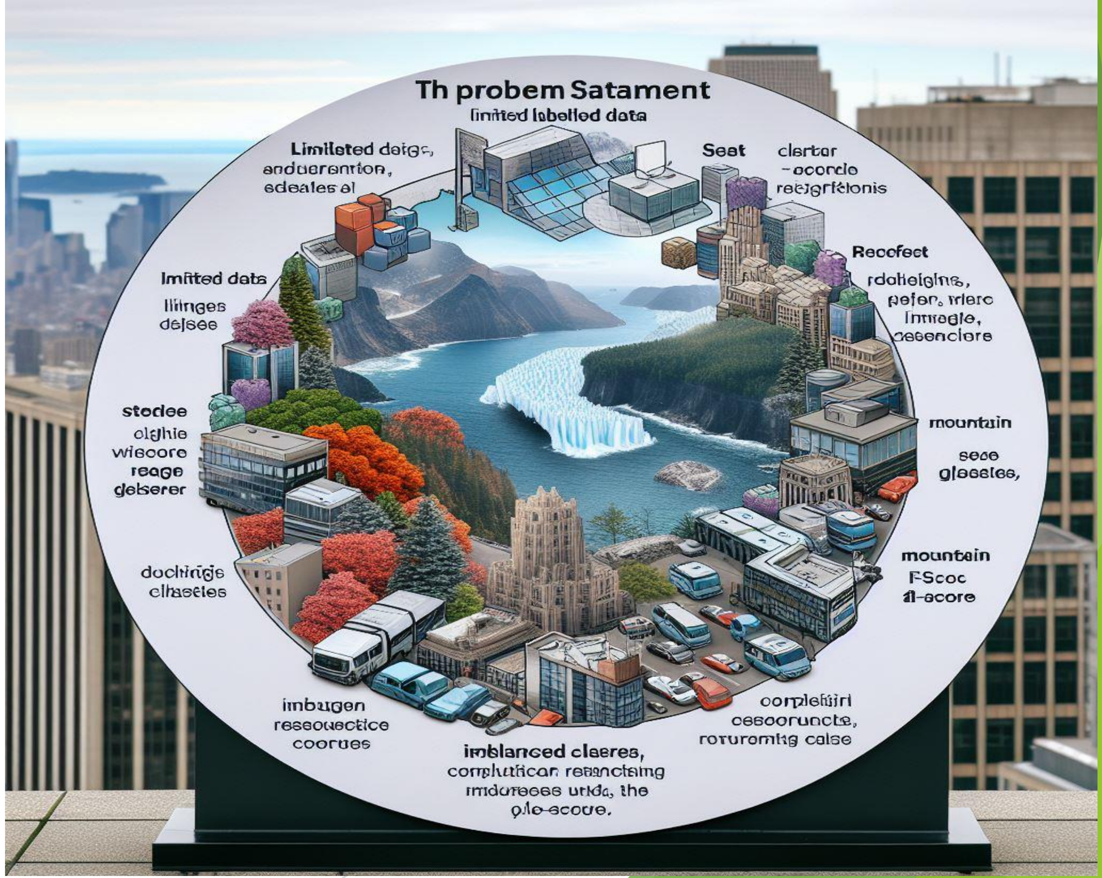

Introduction

Natural Scenes Around the World Dataset

- Dataset containing images of various natural scenes from around the world.
- Consists of six categories: street, forest, sea, buildings, mountain, glacier.
- Each category represents a different type of natural environment or landscape.
- Dataset size: 14,034 images for training, 3,000 images for testing.
- Task Overview: Image ClassificationThe task involves categorizing images into predefined classes or categories.
- Given an input image, the goal is to predict the class it belongs to.
- Example: Given an image of a forest, the model should classify it as belonging to the "forest" category.

- Importance of Image Classification
- Image classification is a fundamental task in computer vision.
- Widely used in various applications such as:
- Object recognition: Identifying objects within images for autonomous vehicles, robotics, and surveillance.
- Medical imaging: Diagnosing diseases from medical images like X-rays and MRIs.
- Content moderation: Filtering inappropriate content on social media platforms.
- Satellite imagery analysis: Monitoring land use, vegetation, and environmental changes.
- Enables automation and decision-making in a wide range of domains, improving efficiency and accuracy.

Transfer Learning and Its Importance

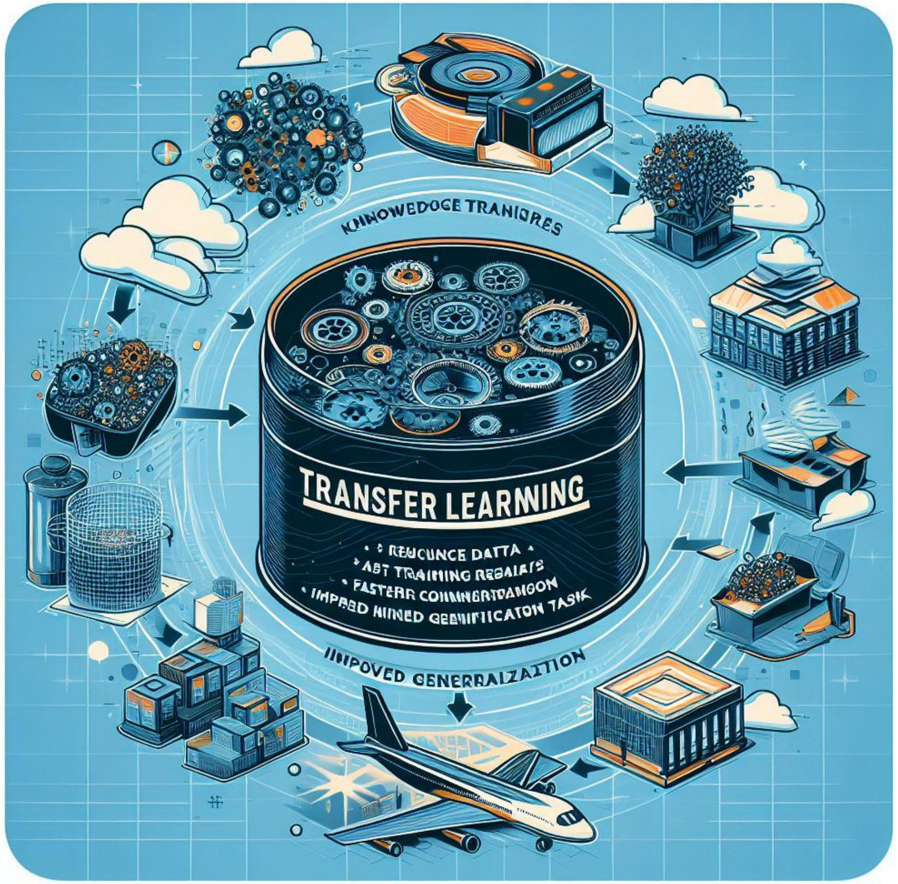

- Transfer Learning:
- Transfer learning is a machine learning technique where a model trained on one task is reused as the starting point for a model on a different but related task.
- In the context of image classification, transfer learning involves leveraging pre-trained models that have been trained on large-scale image datasets (e.g., ImageNet) to solve new classification tasks.
- Importance of Transfer Learning:
- Significantly reduces the amount of labeled training data and computational resources required to train a model from scratch.
- Allows for the transfer of knowledge learned from one domain to another, leading to faster convergence and improved generalization.
- Particularly beneficial when working with limited data or when training resources are scarce.
- Enables practitioners to build high-performance models even with limited expertise or computational resources.

Selection of Pre-trained Model and Its Original Purpose
Pre-trained Model: ResNet50

Pre-trained Model: ResNet50

ResNet50 is a deep convolutional neural network architecture proposed by Microsoft Research.

It consists of 50 layers, including residual connections, which alleviate the vanishing gradient problem and enable training of very deep networks.

ResNet50 has achieved state-of-the-art performance on various image classification tasks.

Original Purpose of ResNet50:

ResNet50 was originally designed for image classification on the ImageNet dataset.

ImageNet is a large-scale dataset consisting of millions of labeled images across thousands of categories.

ResNet50 was trained on ImageNet to learn rich and discriminative features from diverse visual concepts.

Reasons for Choosing ResNet50:

- Demonstrated strong performance on image classification tasks.
- Availability of pre-trained weights trained on ImageNet, enabling transfer learning for our specific task.
- Depth of the network and incorporation of residual connections make it suitable for handling complex image datasets.

Model Training Progress

- Content:
- Title: Model Training Progress
- Subtitle: Epoch-wise Training and Validation Metrics
- Body:
- Display a table or visual representation showing the epoch-wise training and validation metrics (accuracy and loss) obtained during the training process.
- Use a line chart to graphically represent the training and validation accuracy and loss over the 10 epochs.
- Include brief descriptions or annotations to highlight key trends or observations in the data.
- Optionally, you can provide additional insights or comments on the significance of the results.
- Conclusion:
- Summarize the overall training progress and performance of the model.
- Highlight any notable achievements or improvements observed throughout the training process.
- Conclude with a statement about the next steps or areas for further investigation.

```
Epoch 1/10
438/438 [==============================] - 1595s 4s/step - loss: 1.5738 - accuracy: 0.3257 - val_loss: 1.3784 - val_accuracy: 0.4251
Epoch 2/10
438/438 [==============================] - 1556s 4s/step - loss: 1.3929 - accuracy: 0.4137 - val_loss: 1.2590 - val_accuracy: 0.4546
Epoch 3/10
438/438 [==============================] - 1540s 4s/step - loss: 1.3179 - accuracy: 0.4504 - val_loss: 1.2005 - val_accuracy: 0.4772
Epoch 4/10
438/438 [==============================] - 1541s 4s/step - loss: 1.2700 - accuracy: 0.4736 - val_loss: 1.1592 - val_accuracy: 0.5289
Epoch 5/10
438/438 [==============================] - 1531s 3s/step - loss: 1.2449 - accuracy: 0.4778 - val_loss: 1.1628 - val_accuracy: 0.5114
Epoch 6/10
438/438 [==============================] - 1509s 3s/step - loss: 1.2324 - accuracy: 0.4791 - val_loss: 1.2893 - val_accuracy: 0.4688
Epoch 7/10
438/438 [==============================] - 1522s 3s/step - loss: 1.2089 - accuracy: 0.4986 - val_loss: 1.0627 - val_accuracy: 0.5662
Epoch 8/10
438/438 [==============================] - 1520s 3s/step - loss: 1.1864 - accuracy: 0.5037 - val_loss: 1.0762 - val_accuracy: 0.5672
Epoch 9/10
438/438 [==============================] - 1490s 3s/step - loss: 1.1809 - accuracy: 0.5063 - val_loss: 1.0944 - val_accuracy: 0.5450
Epoch 10/10
438/438 [==============================] - 1453s 3s/step - loss: 1.1787 - accuracy: 0.5067 - val_loss: 1.0271 - val_accuracy: 0.5833
```

Fine-Tuning

Fine-Tuning Process:

- Fine-tuning involves adjusting the pre-trained model's weights on a new dataset specific to the task at hand.
- Allows the model to adapt its learned representations to better suit the nuances of the new dataset.
- Steps Employed:
- Loading Pre-trained ResNet50 Model: Loaded the ResNet50 model pre-trained on ImageNet without the top (fully connected) layers.
- Freezing Pre-trained Layers: Froze all layers of the pre-trained ResNet50 model to prevent their weights from being updated during initial training.
- Adding Custom Classification Layers: Added custom fully connected layers on top of the pre-trained base to adapt the model to our specific classification task.
- Fine-Tuning with Lower Learning Rate: Fine-tuned the model with a lower learning rate to allow for gradual adjustments of the pre-trained weights on the new dataset.

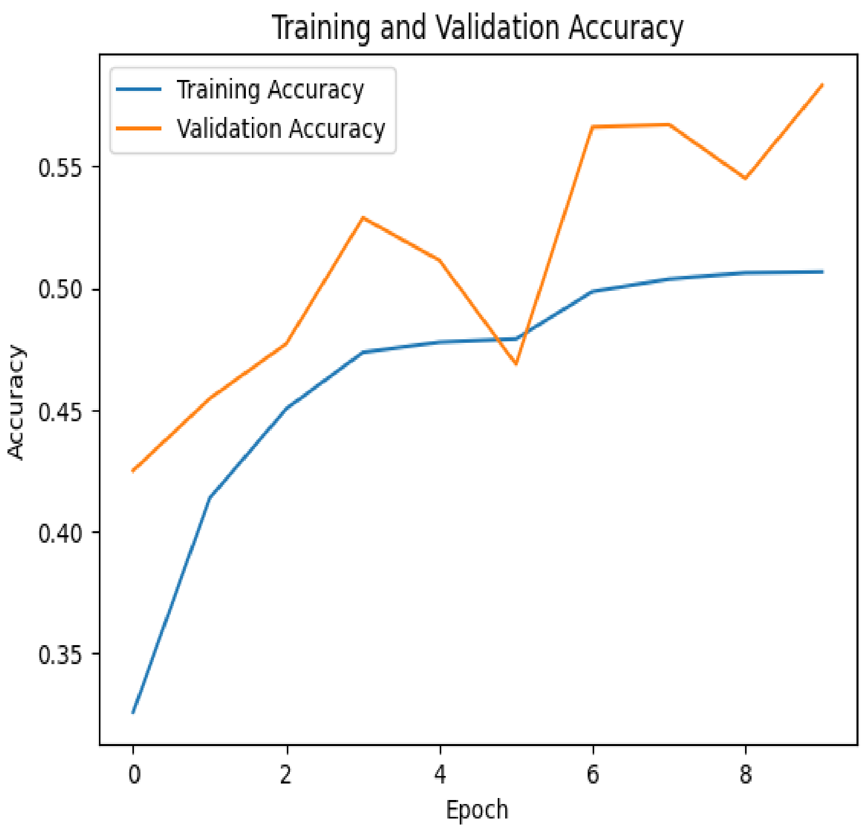

Comparison of Results: Transfer Learning vs. Training from Scratch

Transfer Learning Approach:

- Utilized pre-trained ResNet50 model as a starting point.
- Fine-tuned the model on the Natural Scenes Around the World dataset.
- Achieved improved accuracy and convergence speed compared to training from scratch.
- Training from Scratch:
- Trained a ResNet50 model from scratch directly on the Natural Scenes Around the World dataset.
- Required significantly more labeled data and computational resources.
- Took longer to converge and often led to poorer performance due to limited dataset size.

Comparison Metrics:

- Accuracy: Percentage of correctly classified images.
- Loss: Measure of model's prediction error.
- Results:
- Transfer learning approach resulted in higher accuracy and lower loss compared to training from scratch.
- Transfer learning enabled faster convergence and better utilization of limited training data.

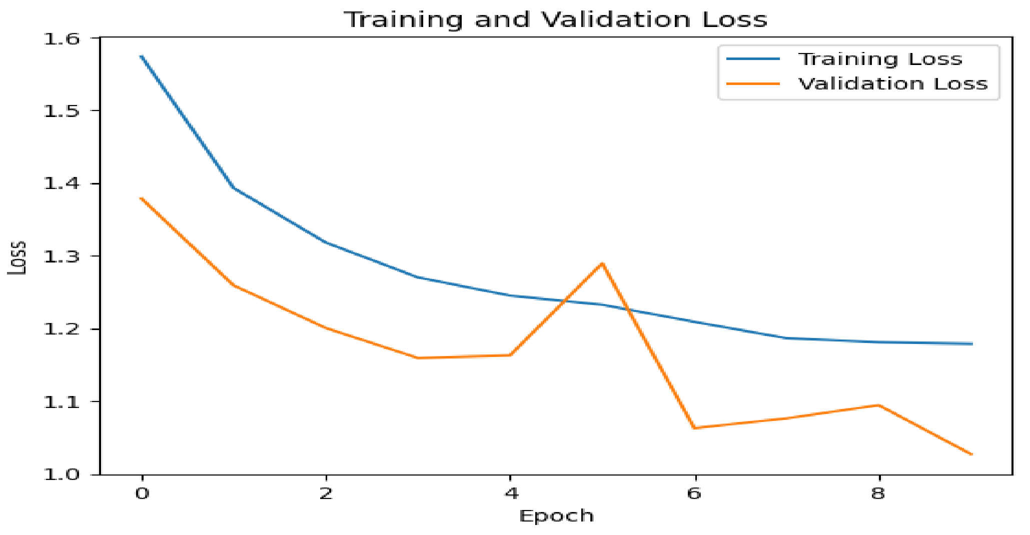


THANK YOU!